


Project Proposal

- **due TODAY @ 11:59pm** 
- one page maximum stating:
 - student names
 - **problem**
 - **data** you will use
 - draft of **proposed solution** (can be updated later)
 - **experiments** you will run (can be updated later)
 - **references** (you can use an additional page for this)
- send me pdf by email (mvasconcelos@eng.ucsd.edu) with:
 - **Subject: Group X Proposal**, where **X** is the group number in this list
 - cc to all group members

This assignment is worth **5% of your class grade**. If you submit the proposal in time and make a serious attempt at addressing the bullet points above, you will get full score. I'm not, at this point, grading projects on their merits. I will look at the proposal and give you some feedback. This will be mostly on issues that I think may become serious obstacles and you need to consider urgently. For example, if I find the problem you propose to be outside the scope of the class, that you may not be able to find data to train the methods you are proposing, etc. Note that if I say "OK", it just means that I see no such problems. It does not mean that you will receive an A just by doing what you proposed. I see these proposals more as a "direction to where the project is going." The projects themselves will be evaluated at the end of the quarter, according to the guidelines published.

1. Hussain, Tanvir; Lewis, Cameron; Villamar, Sandra
2. Dong, Meng; Long, Jianzhi; Wen, Bo; Zhang, Haochen
3. Chen, Yuzhao; Li, Zonghuan; Song, Yuze; Yan, Ge
4. Li, Jiayuan; Xiao, Nan; Yu, Nancy; Zhou, Pei
5. Li, Zheng; Tao, Jianyu; Yang, Fengqi
6. Bian, Xintong; Jiang, Yufan; Wu, Qiyao
7. Chen, Yongxing; Yao, Yanzhi; Zhang, Canwei
8. Nukala, Kishore; Pulleti, Sai; Vaidyula, Srikar
9. Baluja, Michael; Cao, Fangning; Huff, Mikael; Shen, Xuyang
10. Arun, Aditya; Long, Heyang; Peng, Haonan
11. Cowin, Samuel; Liao, Albert; Mandadi, Sumega
12. Jia, Yichen; Jiang, Zhiyun; Li, Zhuofan
13. Dandu, Murali; Daru, Srinivas; Pamidi, Sri
14. He, Bolin; Huang, Yen-Ting; Wang, Shi; Wang, Tzu-Kao
15. Chen, Luobin; Feng, Ruining; Wu, Ximei; Xu, Haoran
16. Chen, Rex; Liang, Youwei; Zheng, Xinran
17. Aguilar, Matthew; Millhiser, Jacob; O'Boyle, John; Sharpless, Will
18. Wang, Haoyu; Wang, Jiawei; Zhang, Yuwei
19. Chen, Yinbo; Di, Zonglin; Mu, Jiteng
20. Chowdhury, Debalina; He, Scott; Ye, Yiheng
21. Lin, Wei-Ru; Ru, Liyang; Zhang, Shaohua
22. Bhavsar, Shivad; Blazej, Christopher; Bu, Yinyan; Liu, Haozhe
23. Chen, Claire; Hsieh, Chia-Wei; Lin, Jui-Yu; Tsai, Ya-Chen
24. Cheng, Yu; Yu, Zhaowei; Zaidi, Ali
25. Assadi, Parsa; Brugere, Tristan; Pathak, Nikhil; Zou, Yuxin
26. Candassamy, Gokulakrishnan; Dixit, Rajeev; Huang, Joyce
27. Kok, Hong; Wang, Jacky; Yan, Yijia; Yuan, Zhouyuan
28. Luan, Zeting; Yang, Zheng
29. Cuawenberghs, Kalyani; Mojtahed, Hamed

ECE 271B – Winter 2022

Neural Networks (cont.)

Disclaimer:


This class will be recorded
and made available to students asynchronously.

Manuela Vasconcelos
ECE Department, UCSD

Neural Network


- the MLP as function approximation

even with just 2 layers,
it is possible
to approximate
complicated functions!



Two Modes of Operation

- normal mode, after training: feedforward

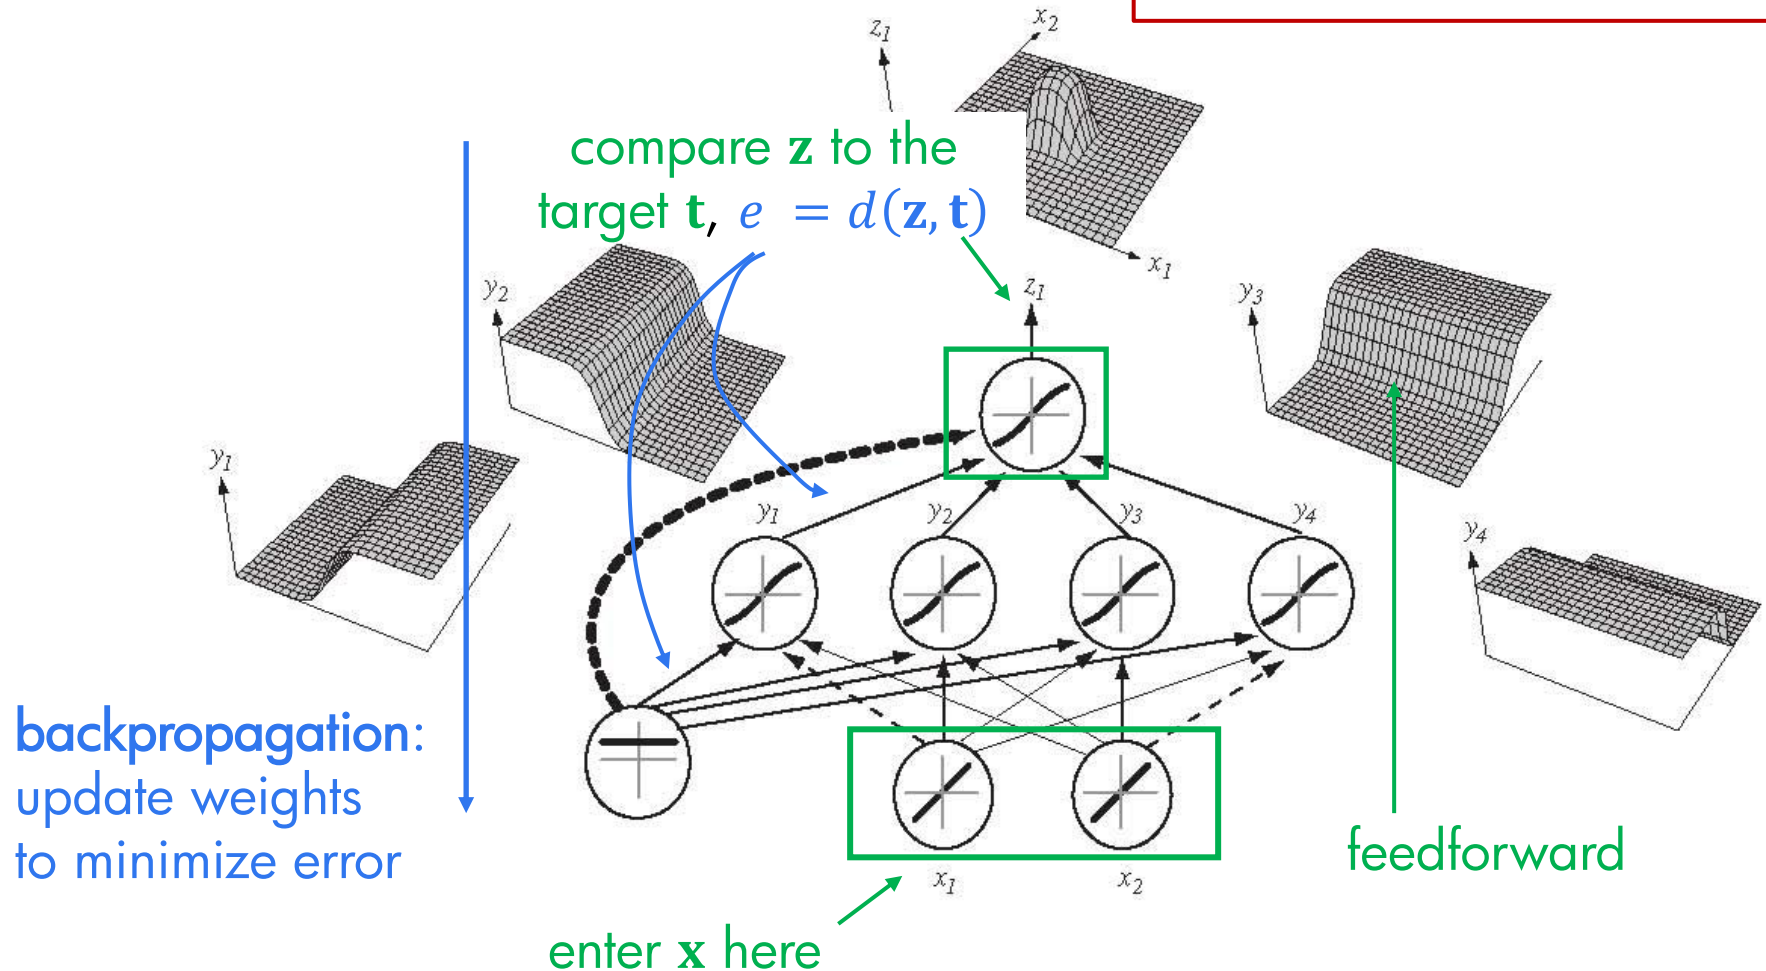


Two Modes of Operation

► training mode: backpropagation

iterative:

- propagate **x** feedforward
- compute **error e**
- **backpropagate** to adjust weights



Backpropagation

- ▶ is just gradient descent
- ▶ at the end of the day, the **output \mathbf{z}** is just a big function of
 - **input vector \mathbf{x}**
 - **weights**, which we can be represent by a “big” vector **\mathbf{W}**
 - e.g.

$$\mathbf{z} = s \left[\sum_j v_j s \left(\sum_i w_{ji} x_i \right) \right] = \mathbf{z}(\mathbf{x}; \mathbf{W}) \quad \text{with} \quad \mathbf{W} = (\mathbf{v}, \mathbf{w})$$

- ▶ **objective:** given a dataset $\mathcal{D} = \{(\mathbf{x}_1, \mathbf{t}_1), \dots, (\mathbf{x}_n, \mathbf{t}_n)\}$, determine

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} J(\mathbf{W})$$

$$J(\mathbf{W}) = \sum_{i=1}^n L(\mathbf{t}_i, \mathbf{z}(\mathbf{x}_i; \mathbf{W}))$$


$$L(\mathbf{t}, \mathbf{z}) = \frac{1}{2} \sum_k [t_k - z_k]^2$$

squared-error with the multi-label classifier

Feature Transformation

- MLP can be seen as:

non-linear feature transformation + linear discriminant



Feature Transformation

► feature transformation



- searches for the **space** where the **patterns become separable**

► Q: is **separability** always possible?

► A: **not** really, depends on the **number of units**

► in practice

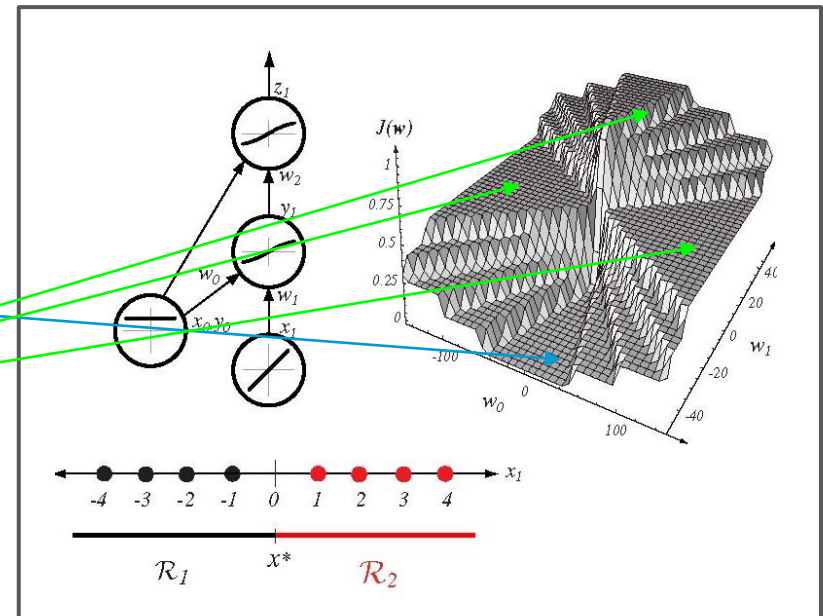
- art—form
- trial and error



Other Problems

► the optimization surface (cost) can be quite nasty

- cost has many “plateaus”
- global optimal solution has no error
 - but gradient frequently close to zero
 - slow progress



► how do we set the learning rate η ?

- if too small or too big, we will need various iterations
- could even diverge

Structural Risk Minimization

- ▶ what about **complexity penalties**, **overfitting**, and **all that**?
- ▶ solve

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \sum_{i=1}^n L(\mathbf{t}_i, \mathbf{z}(\mathbf{x}_i; \mathbf{W})) \quad \text{subject to} \quad \mathbf{W}^T \mathbf{W} < \lambda$$

- ▶ we will see that this is equivalent to

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \left\{ \sum_{i=1}^n L(\mathbf{t}_i, \mathbf{z}(\mathbf{x}_i; \mathbf{W})) + \frac{2\varepsilon}{\eta} \mathbf{W}^T \mathbf{W} \right\}$$

- ▶ re—working out backpropagation, this can be done by “**shrinking**”
 - after each weight update, do $\mathbf{W}^{new} = \mathbf{W}^{old}(1 - \varepsilon)$, $0 < \varepsilon < 1$
 - this is known as “**weight decay**” and **penalizes complex models**

In Summary

- ▶ this works, but requires **tuning ε**
- ▶ the **cost surface** is **nasty**
- ▶ one needs to try **different architectures**
- ▶ hence, training can be **painfully slow**
 - “**weeks**” is quite common
 - a good neural network may take **years** to train
- ▶ however, when you are finished it **tends to work well**
- ▶ examples
 - the Rowley and Kanade face detector
 - the LeCun digit recognizer (see <http://yann.lecun.com/exdb/lenet/index.html>)

The Last Five–Ten Years

- ▶ what we have seen so far was the **state of the art** in the 1990's
- ▶ over the last 5–10 years, neural networks have become **dominant again**
- ▶ you surely have heard all the talk about “**deep learning**” and how it is **revolutionizing AI**
- ▶ what has happened?
well, a **combination** of

- a few (small) advances in neural network architectures
- large advances in
 - • implementation
 - availability of data (what is now called “Big Data”)

Architecture

- ▶ there have been some **advances in NN architectures** (although some of these were already used in the 80s)
- ▶ a **popular architecture** is the **Convolutional NN (CNN)**
 - commonly used in applications involving **signals** (images, speech, audio)
 - inspired by classic linear systems theory
- ▶ consider a **linear filter**
 - for simplicity, we consider a 1D signal $x[n]$
 - to which we apply a **filter** of impulse response $h[n]$ (e.g. to remove noise)
 - mathematically, the **filter output** is given by the **convolution** of x and h

$$y[n] = \sum_k x[k]h[n - k]$$

Architecture

► linear filter

- the **convolution** of x and h


$$y[n] = \sum_k x[k]h[n - k]$$

can be written as

$$y[n] = \sum_k x[k]g_n[k] = \langle x[k], g_n[k] \rangle \quad g_n[k] = h[n - k]$$


i.e., each entry of $y[n]$ is the **dot-product** of $x[n]$ with a shifted and inverted replica of the impulse response $h[n]$

- e.g., let






Architecture

► convolution of x and h



$$y[n] = \sum_k x[k]g_n[k] = \langle x[k], g_n[k] \rangle \quad g_n[k] = h[n - k]$$




and so on...

Architecture

► convolution of x and h

$$y[n] = \sum_k x[k]g_n[k] = \langle x[k], g_n[k] \rangle \quad g_n[k] = h[n - k]$$

- this can be computed as



Note




for simplicity of the picture,
 $y[1]$, $y[3]$ are not represented

► note that this is a **neural network** with two properties

- sparse connectivity: each unit only has a few non-zero weights
- shared weights: the weights of all units are equal

Architecture


- ▶ a network with such layers is called **convolutional**
- ▶ **Convolutional NNs (CNNs)** are very suitable for **signal processing** type of applications (vision, speech)
- ▶ besides this, it has been found that the training problems can be alleviated by using **different non-linearities**
- ▶ instead of the traditional **hyperbolic tangent** or the **sigmoid**
- ▶ modern networks frequently use **rectified linear units**



Architecture

► Rectified Linear Unit (ReLU)

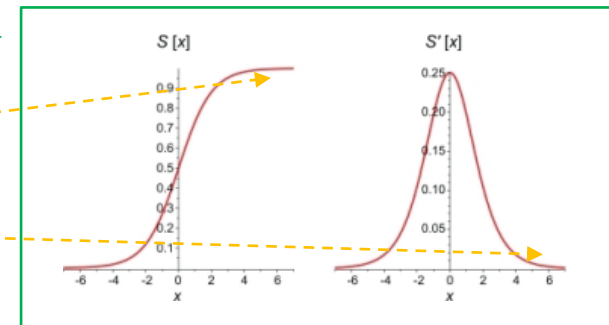
$$s(x) = \max(x, 0)$$



$$\frac{\partial L}{\partial w_{ji}} = -\delta_j y_i$$
$$\delta_j = [\sum_k \delta_k w_{kj}] s'[g_j] \quad \text{if } j \text{ is hidden}$$
$$\delta_j = (t_j - z_j) s'[u_j] \quad \text{if } j \text{ is output}$$

► recall: the gradient of backpropagation depends only on the **derivative of the non-linearity**

- for **sigmoid** or **hyperbolic tan**, this is mostly zero
- a **large response** creates **zero derivative**
- for networks that have many layers, the **gradient vanishes**
- since the **ReLU** has **constant derivative**, it reduces this problem
- learning tends to converge much faster



Architecture

- ▶ another factor that can significantly **speed up** the training of neural networks is the **cost function**

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} J(\mathbf{W})$$

$$J(\mathbf{W}) = \sum_{i=1}^n L(\mathbf{t}_i, \mathbf{z}(\mathbf{x}_i; \mathbf{W}))$$

$$L(\mathbf{t}, \mathbf{z}) = \frac{1}{2} \sum_k [t_k - z_k]^2$$

- ▶ above, we have used the **square of the error**

$$[t - z(\mathbf{x}; \mathbf{W})]^2$$

- ▶ a more popular choice nowadays is the **cross-entropy loss**

$$-t \log z(\mathbf{x}; \mathbf{W}) + (1 - t) \log(1 - z(\mathbf{x}; \mathbf{W}))$$

- ▶ for networks with **sigmoid units**, this again **eliminates** the dependence of the **gradient** on the derivative of the sigmoid

Computing the Gradient of L

$$\begin{aligned}
 \frac{\partial L}{\partial y_j} &= \frac{\partial}{\partial y_j} \left[\sum_k t_k \log z_k + (1 - t_k) \log(1 - z_k) \right] \\
 &= - \sum_k \left(\frac{t_k}{z_k} - \frac{1 - t_k}{1 - z_k} \right) \frac{\partial z_k}{\partial y_j} \\
 &= - \sum_k \frac{t_k - z_k}{z_k(1 - z_k)} \frac{\partial z_k}{\partial u_k} \frac{\partial u_k}{\partial y_j} \\
 &= - \sum_k \frac{t_k - z_k}{z_k(1 - z_k)} s'[u_k] w_{kj}
 \end{aligned}$$


- for the **sigmoid non-linearity**, it can be shown that

$$s'[u_k] = z_k(1 - z_k)$$

and

$$\frac{\partial L}{\partial y_j} = - \sum_k (t_k - z_k) w_{kj}$$

no longer depends on $s'[u_k]$



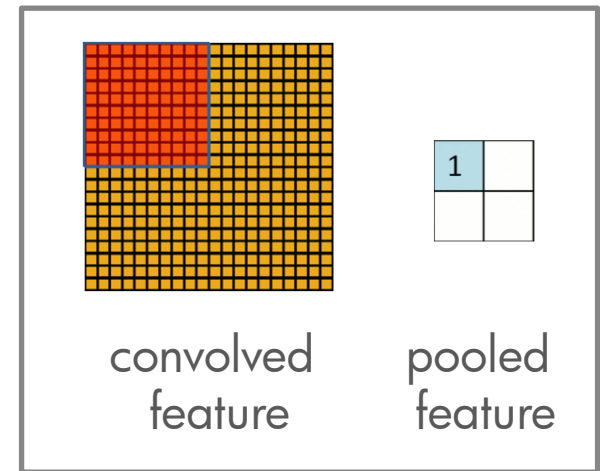
$$z_k = s[u_k]$$

$$u_k = \sum_j w_{kj} y_j$$

Architecture

► CNN usually also perform **pooling**

- this consists of **mapping a region of the convolution output** (after the non-linearity) into a **single response**
- in the example, 10×10 units are mapped into one value
- **pooling operators** are typically the **average** or the **maximum**
 - consider **maximum pooling**: as long as the largest response among the orange units stays the same, the pooled output is the same
- this has two benefits
 - **dimensionality reduction**: from 40×40 units to 4
 - **invariance**: the pooled response is invariant to small shifts or changes of scale (size) of the input



Architecture


► deep learning

- deep learning networks are **networks with many layers**
- modern deep learning networks, usually involve a combination of **convolutional** and **fully-connected layers**, **ReLUs**, and **pooling**
- this started with the **AlexNet**
 - 5 convolutional layers
 - some with max pooling
 - 3 fully-connected (conventional) layers
 - ReLU non-linearities
 - around 650,000 units, 60 million parameters (weights)
 - created a big buzz by significantly outperforming the previous best results in the problem of object recognition

AlexNet

Krizhevsky, Alex; Sutskever, Ilya; Hinton, Geoffrey;
ImageNet Classification with Deep Convolutional
Neural Networks. *NIPS* 25, 1097–1105, 2012.

The class with the highest likelihood is the one the DNN selects



When AlexNet is processing an image, this is what is happening at each layer.

NOTE:


Today, there are also GoogleNet, ResNet (Microsoft) - 152 layers, VGG (Oxford), etc.

Implementation

► complexity


- as you might imagine, a network like this is not easy to train or use
- an important development has been the introduction of **GPUs (Graphical Processing Units)**
 - these are the processors that come with your graphics board
 - they are **specialized processors** that can **perform convolutions much faster** than traditional CPUs
 - all modern CNNs are implemented on GPUs, using publically available software packages
 - typically, you do not train a CNN but adapt an existing one (e.g., AlexNet) to your problem
 - this consists of **initializing** your network with the existing one and **running a few iterations** of backpropagation on your data
 - this is called **fine-tuning**

AlexNet



Fine-Tuning

- ▶ MLP can be seen as: **non-linear feature transformation** + **linear classifier**
- ▶ **Fine-tuning** basic idea is
 - keep the **feature transformation**
 - replace the **linear classifier** by one **suitable** for the new problem
- ▶ **Procedure:**
 - replace the **last (output) layer** by one with as **many output units** as the **number of classes** in your problem
 - keep the rest of the CNN exactly the **same**
 - **run backpropagation** on the new network
 - **limit the number of iterations** so that it does not **overfit** (if your training set is small)



Big Data

► complexity also affects learning

- how many examples do you need to learn 60 million parameters?
- training of these networks has only become possible with the **advent of large datasets**
- the most popular one is **ImageNet**, which (currently) contains 14 million images of thousands of classes


*Deng, J.; Dong, W.; Socher, R.; Li, L.-J.; Li, K.; Fei-Fei, L.
ImageNet: A Large-Scale Hierarchical Image Database. CVPR, 2009.*



- again, you typically do **not** collect a dataset of this magnitude, but **fine-tune** the network trained on ImageNet using your (smaller) training set

Examples

- ▶ the joint results of **architecture advances**, **deep learning**, **GPUs**, and **big data** have been staggering
 - today, we have vision and speech systems that could not be imagined 5 years ago
 - tasks like object recognition now seem to be solvable



*Cai, Zhaowei; Vasconcelos, Nuno;
Cascade R-CNN: Delving into High
Quality Object Detection. CVPR 2018.*

<https://www.youtube.com/watch?v=l9kNhXfNnHs>

- in a decade, deep learning will be in your car, phone, TV, etc.